



Ablation Study — E · KL Temperature β (GPT-2, Anthropic/hh-rlhf)


Training Loss




Train Reward Accuracy




Train Reward Margin




Clip Fraction




Eval Reward Accuracy



Eval Reward Margin



Eval Token-mean KL



Final Eval Accuracy

