



Ablation Study — $E \cdot \text{KL Temperature } \beta$ (GPT-2, Anthropic/hh-rlhf)


Training Loss




Train Reward Accuracy




Train Reward Margin




Clip Fraction




Eval Reward Accuracy



Eval Reward Margin



Eval Token-mean KL



Final Eval Accuracy

